

The AI assistant: what each role can do

`/chat` is an HR assistant that reads real data and takes real actions. It answers handbook questions with citations, reports balances, books time off, and, for the right roles, approves requests and surfaces retention signals.

The interesting part is not that it does those things. It is that it cannot do them for the wrong person. This guide demonstrates that per role, with screenshots taken from a running instance signed in as the demo account named in each caption.

How the assistant is limited

The assistant works through tools, and a role's tools are chosen before the conversation starts. A tool the role may not use is never given to the model, so it is not something the model declines. It is something the model has no way to see or attempt, which is also why a prompt-injection attempt cannot reach it.

The app renders this matrix from the same function that builds the toolset, so the table below is the running configuration rather than a description of it. It lives at `/settings/ai-tools`, visible to Super Admin.

Tool	Employee	Manager	HR Admin	Super Admin
getCurrentDateTime The current date, time, and weekday (utility).	✓	✓	✓	✓
getDateInfo Weekday / weekend status for a given date (utility).	✓	✓	✓	✓
businessDaysBetween Working-day count between two dates (utility).	✓	✓	✓	✓
endConversation End the conversation and lock the chat — used only to decline abuse or unsafe requests.	✓	✓	✓	✓
searchHandbook Search the employee handbook (RAG) and cite sections.	✓	✓	✓	✓
getEmployeeDirectory List employees you're allowed to see.	✓	✓	✓	✓
getLeaveBalances Your own time-off balances.	✓	✓	✓	✓
requestTimeOff Submit a time-off request for you.	✓	✓	✓	✓
listPendingApprovals List requests awaiting your approval.	✗	✓	✓	✓
approveLeave Approve or reject a pending request.	✗	✓	✓	✓
getPayslip A payslip — your own, or anyone visible with elevated rights.	✓	✓	✓	✓
predictDepartures Rank who is most at risk of leaving (managers: own team, anonymized; HR: full detail).	✗	✓	✓	✓
getEngagementRisk Find who is at risk of burnout/disengagement (managers: own team; HR: company). Names never exposed to the model.	✗	✓	✓	✓

Read the columns and the shape of the thing becomes clear:


- An Employee gets 9 of the 13 tools. The four they lack are approvals (two) and the predictive and engagement tools.
- Manager, HR Admin and Super Admin get the **same 13 tool names**. `admin:settings` unlocks no extra tools at all, so a Super Admin's assistant is identical to an HR Admin's.

That last point matters for reading the rest of this guide. From Manager upward, the difference between roles is not which tools they have. It is what those tools return, and what shape they take.

One question, three roles

The cleanest demonstration in the app. The same question, asked by three people.

Employee. `predictDepartures` was never injected, so there is no tool call, no card, and no refusal. The model answers in prose and describes its actual limits, then offers what it can do instead.



There is no "access denied" card here, and that is deliberate rather than an omission. An out-of-scope request is not an error, so the assistant treats it as an ordinary question it happens not to be able to answer.

Manager. The tool runs, scoped to their team. The card names departments, never people, and says so on its face.

 **HARI**
HUMAN RESOURCES · AI

MANAGER WORKSPACE

Dashboard

AI Assistant

Directory

My Team

Time Off

Onboarding

Documents

Knowledge Base

KE

Karim El Idrissi
Manager

HARI / AI Assistant

Q Search...



Manager

AI Assistant

Manager

+ New chat

openrouter-auto

Who is at risk of leaving?



 Thought process

I'll check which of your team members are at highest risk of leaving.

Analyzing departure risk predictDepartures

Departure risk

Team

0 of 7 high-risk

Aggregate team view — individuals aren't named.

#1 Design Moderate 58

Engagement score Salary evolution Time since last review

#2 Sales Moderate 49

Burnout (unused PTO) Role stagnation Department turnover


#3 IT Low 37

Burnout (unused PTO) Engagement score Seniority

Ask anything... (Shift+Enter for newline)



HR Admin. The same tool, the same card component, company scope, named individuals.



Side by side:


	Employee	Manager	HR Admin
Tool available	no	yes	yes
Scope	none	own team (7 people)	company (59 people)
Top row reads	prose answer	#1 Design	#1 Asmae Tazi
Names	none	never	yes
Salary	never	never	never

The anonymization is not a display trick. The manager's copy of the data never contains a name or an employee id, so the model cannot cross-reference the id against the directory to re-identify anyone. It only ever received a department, a score, a band, and factor keys.


One honest limitation, since this is a demonstration and not a sales page: department-level anonymity is only as strong as the department is large. A manager whose team has exactly one designer can infer who #1 Design is. The mechanism removes the identifier; it does not promise k-anonymity.

Where each role starts

The suggested prompts are filtered by role, so the assistant never offers something its tools cannot deliver. The first two suggestions are shared. The last two are not.



Employee: Show me my profile, Show me my latest payslip.



Manager: Show me the team directory, What time-off requests are awaiting my approval?

The role badge next to the assistant's name is the quickest way to read any screenshot in this guide.

Handbook questions

Every role can ask the handbook. The assistant retrieves sections, cites them by number, and links each one to its exact anchor.

The screenshot displays the HARI AI Assistant interface. On the left is a dark sidebar with the HARI logo and a menu for 'EMPLOYEE WORKSPACE' including Dashboard, AI Assistant (selected), Directory, Time Off, Onboarding, Documents, and Knowledge Base. The main area shows a chat window titled 'HARI / AI Assistant' with a search bar and a user dropdown set to 'Employee'. A blue button asks 'What is our parental leave policy?'. The assistant's response shows a 'Thought process' of 'Searching the handbook' followed by 'Handbook results for "parental leave policy" (4)'. The final answer states: 'Parental leave includes 14 weeks of maternity leave, of which 7 weeks must be taken after the birth. Paternity leave is 3 working days within the month following the birth. Employees must contact HR at least 30 days before their intended start date [1]'. A citation '[1]' is present at the end of the text. The bottom of the interface has a text input field 'Ask anything... (Shift+Enter for newline)' and a send button.

The [1] in the answer is a real link into the Knowledge Base. The model does not compose these URLs; it emits a number, and the link is built server-side from the retrieved record. A citation therefore cannot point at a document that was not retrieved, or at one the asker may not read.

Opening the panel shows what was actually retrieved and how well each section matched.

The screenshot displays the HARI AI Assistant interface. On the left is a dark sidebar with the HARI logo and 'HUMAN RESOURCES · AI'. Below the logo is the 'EMPLOYEE WORKSPACE' section with links to Dashboard, AI Assistant (highlighted), Directory, Time Off, Onboarding, Documents, and Knowledge Base. At the bottom of the sidebar is a user profile for 'Imane Chaibi' (Employee). The main panel has a header with 'HARI / AI Assistant', a search bar, and a user dropdown set to 'Employee'. Below the header, there's a '+ New chat' button and a dropdown menu showing 'openrouter-auto'. The chat area shows a question: 'What is our parental leave policy?'. Below the question, a 'Thought process' section indicates 'Searching the handbook' with a 'searchHandbook' link. The results are titled 'Handbook results for "parental leave policy" (4)' and list three items:

- [1] **Parental Leave** (56% match)
Leave & Time Off
Maternity leave is 14 weeks, including 7 weeks that must be taken after the birth. Paternity leave is 3 working days, taken within the month following the birth. Contact HR at least 30 days before your intended start date.
- [2] **Paid Time Off (Vacation)** (42% match)
Leave & Time Off
Full-time employees accrue 1.5 working days of paid leave per month of service — 18 working days per year for a full year — rising with seniority up to a maximum of 30 days. Vacation must be requested in advance...
- [3] **Sick Leave** (32% match)
Leave & Time Off
Any sick absence must be justified by a medical certificate sent to HR within 48 hours; partial salary is maintained in line with CNSS rules, subject to eligibility. Sick leave can be used for personal illness or to care...

At the bottom of the chat area is an input field with the placeholder 'Ask anything... (Shift+Enter for newline)' and a send button.

Retrieval is scoped to the asker's tier, exactly as the Knowledge Base reader is. The same question from a Manager can surface a section an Employee will never be shown. See [knowledge-base.md](#) for that gate.

Self-service, and a schema that cannot express the wrong question

Asking for your own payslip works at every role.

The screenshot shows the HARI AI Assistant interface. On the left is a dark sidebar with the HARI logo and a menu: EMPLOYEE WORKSPACE, Dashboard, AI Assistant (selected), Directory, Time Off, Onboarding, Documents, and Knowledge Base. At the bottom of the sidebar is a user profile for Imane Chaibi, Employee. The main area has a header with 'HARI / AI Assistant', a search bar, and a dropdown for 'Employee'. Below the header is a 'New chat' button and a dropdown for 'openrouter-auto'. A blue button 'Show me my latest payslip' is visible. The chat history shows a 'Thinking...' message with the following reasoning: 'We need to respond with tool call to getPayslip, as the user wants to see latest payslip. According to policy: For READS (payslip), we must call the tool in the same turn; we cannot just describe. So we need to call getPayslip with no parameters (the schema shows no parameters). We'll do that.' This is followed by a 'Retrieving payslip' message and a table for Imane Chaibi's payslip details. Another 'Thinking...' message follows, stating: 'Now we need to respond with a concise summary: show the payslip details: gross, tax, net. Provide in MAD. Should not repeat raw JSON but give summary. Also'. At the bottom is a text input field with the placeholder 'Ask anything... (Shift+Enter for newline)' and a send button.

Imane Chaibi	
Most recent month	
Gross	MAD 15,000
Tax (22%)	-MAD 3,300
Net	MAD 11,700

The reasoning panel in that screenshot is worth reading closely. The model says it will call the tool "with no parameters (the schema shows no parameters)". That is the mechanism working: for a role without `payslip:read:any`, the payslip tool is defined with an empty input schema. There is no employee field to fill in.

So "show me Sara's payslip" is not refused for an Employee. It is inexpressible. The request cannot be represented as a valid tool call, which is a stronger guarantee than a rule the model is asked to follow. `getPayslip` is the only tool whose schema changes shape by role; everywhere else the schema is constant and the scoping happens server-side.

The same tool, different data

The directory tool is available to every role, and returns something different for each. Both screenshots below contain Yassine Bouzid.

The screenshot displays the HARI AI Assistant interface. On the left is a dark sidebar with the HARI logo and a menu for 'MANAGER WORKSPACE' including Dashboard, AI Assistant (highlighted), Directory, My Team, Time Off, Onboarding, Documents, and Knowledge Base. The main area is titled 'HARI / AI Assistant' and shows a search bar and a dropdown menu set to 'Manager'. Below this, a grid of employee cards is displayed. Each card includes a circular icon with initials, the employee's name and title, their department and role, location, and email. The employees listed are: Hamza Tahiri (Marketing, Employee, Tanger), Imane Chraïbi (IT, Employee, Tétouan), Karim El Idrissi (IT, Manager, Tétouan) - marked as 'You', Leïla Haddadi (Finance, Employee, Rabat), Nour Ouazzani (IT, Employee, Tétouan), Omar Alaoui (Infrastructure, Employee, Rabat), Sara Amrani (Design, Employee, Casablanca), and Yassine Bouzid (Sales, Employee, Casablanca). At the bottom, there is a text input field with the placeholder 'Ask anything... (Shift+Enter for newline)' and a submit button.

Initials	Name	Title	Department	Role	Location	Email
HT	Hamza Tahiri	Chargé Marketing	Marketing	Employee	Tanger	h.tahiri@hari.ma
IC	Imane Chraïbi	Ingénieure Logiciel Full Stack	IT	Employee	Tétouan	collaborateur@hari.ma
KE	Karim El Idrissi	Manager Technique	IT	Manager	Tétouan	manager@hari.ma
LH	Leïla Haddadi	Contrôleuse de Gestion	Finance	Employee	Rabat	l.haddadi@hari.ma
NO	Nour Ouazzani	Stagiaire Data	IT	Employee	Tétouan	n.ouazzani@hari.ma
OA	Omar Alaoui	DevOps Engineer	Infrastructure	Employee	Rabat	o.alaoui@hari.ma
SA	Sara Amrani	UX/UI Designer	Design	Employee	Casablanca	s.amrani@hari.ma
YB	Yassine Bouzid	Commercial Grands Comptes	Sales	Employee	Casablanca	y.bouzid@hari.ma

Manager: own team. Name, title, department, location, email, and no pay.

The screenshot displays the HARI AI Assistant interface. On the left is a dark sidebar with navigation links: Dashboard, AI Assistant (highlighted), Directory, My Team, HR Activity, HR Analytics, Time Off, Onboarding, Offboarding, Documents, Knowledge Base, Alerts, and Audit Log. At the bottom of the sidebar is a user profile for Nadia Benali, HR Admin. The main content area is titled 'HARI / AI Assistant' and features a search bar and a '+ New chat' button. Below these are several employee cards, each showing a profile picture, name, title, department, location, email, and salary. The cards are for YB Tassine Bouzid (Sales, Employee, Casablanca, MAD 175,000 / yr), YB Younes Benini (Operations, Manager, Casablanca, MAD 200,000 / yr), YC Younes Chraïbi (Design, Employee, Fès, MAD 158,000 / yr), YB Youssef Benali (Engineering, Manager, Casablanca, MAD 245,000 / yr), YB Youssef Berrada (Finance, Employee, Marrakech, MAD 186,000 / yr), YT Youssef Tazi (DSI, Super Admin, Distanciel, MAD 350,000 / yr), and ZH Zakaria Hakimi (Design, Employee, Marrakech, MAD 182,000 / yr). At the bottom of the main area is a text input field with the placeholder 'Ask anything... (Shift+Enter for newline)' and a submit button.

Employee ID	Name	Title	Department	Location	Email	Salary
YB	Tassine Bouzid	Commercial Grands Comptes	Sales	Casablanca	y.bouzid@hari.ma	MAD 175,000 / yr
YB	Younes Benini	Operations Manager	Operations	Casablanca	lead.operations@hari.ma	MAD 200,000 / yr
YC	Younes Chraïbi	UX Designer	Design	Fès	design.16@hari.ma	MAD 158,000 / yr
YB	Youssef Benali	Engineering Manager	Engineering	Casablanca	lead.engineering@hari.ma	MAD 245,000 / yr
YB	Youssef Berrada	Accountant	Finance	Marrakech	finance.32@hari.ma	MAD 186,000 / yr
YT	Youssef Tazi	Directeur des Systèmes d'Information	DSI	Distanciel	admin@hari.ma	MAD 350,000 / yr
ZH	Zakaria Hakimi	UI Designer	Design	Marrakech	design.14@hari.ma	MAD 182,000 / yr


HR Admin: the whole company, and MAD 175,000 / yr on that same person.

Managers are pay-blind. Salary requires `salary:read:all`, which only HR Admin and Super Admin hold, and the field is stripped in the shared data layer rather than hidden in the card. The `/directory` page and the assistant call the same function and receive the same object, so the two surfaces cannot disagree.

Guardrails

A regular expression check runs over every message in the turn before any model call. It is deliberately narrow, anchored to override verbs, so that ordinary questions ("what are your instructions for booking leave?") pass untouched.

When it fires, no model call happens at all. The conversation is locked, the composer is replaced, and the turn is recorded and raised as an alert to HR and admin roles.



This is the outermost of three layers, and the least important of them. The first and strongest is the one from the top of this guide: a role's tools are chosen before the conversation starts, so the injection above had nothing to reach even if it had passed. The third is a tool the model can call itself to end a conversation that has become abusive or unsafe.

What is recorded

Every turn writes an event: role, model, token counts, latency, finish reason, tool name, and outcome codes. No prompts, no responses, no names, no salary. The observability trail is metadata only, by design and for compliance, and the same restraint applies to the predictive tools above, which is why a manager's card carries factor keys rather than sentences about a named person.

Reproducing this

Sign in at `/login` with any demo account (all use `password123`) and ask the same question from each:

Account	Role	"Who is at risk of leaving?"
<code>collaborateur@hari.ma</code>	Employee	prose answer, no tool call
<code>manager@hari.ma</code>	Manager	anonymized card, team scope
<code>rh@hari.ma</code>	HR Admin	named card, company scope
<code>admin@hari.ma</code>	Super Admin	identical to HR Admin

`/settings/ai-tools` renders the live matrix as Super Admin. For the tool contract, the refusal and error result shapes, and the checklist for adding a tool, see [docs/architecture/authorization-invariants.md](#). The permission definitions are in `src/lib/rbac.ts`, and the tool catalogue in `src/lib/ai/tools.ts`.

A note on the model

The screenshots show `openrouter-auto`, the default. It routes to whichever free model is currently available, so the exact wording of any answer will differ from run to run; the cards, the scoping, and the citations will not. Model choice affects phrasing, never permissions.

If a demonstration has to be reliable on the day, set `OPENAI_API_KEY` and pick one of the GPT models marked "(paid)" in the picker. They bill a real account, which is why they are not the default, but they cannot be retired from under you. The free model this project originally pinned was withdrawn by its provider, which is the failure that option exists to avoid.